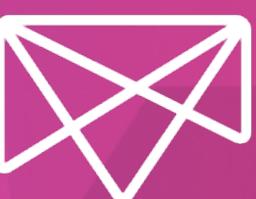


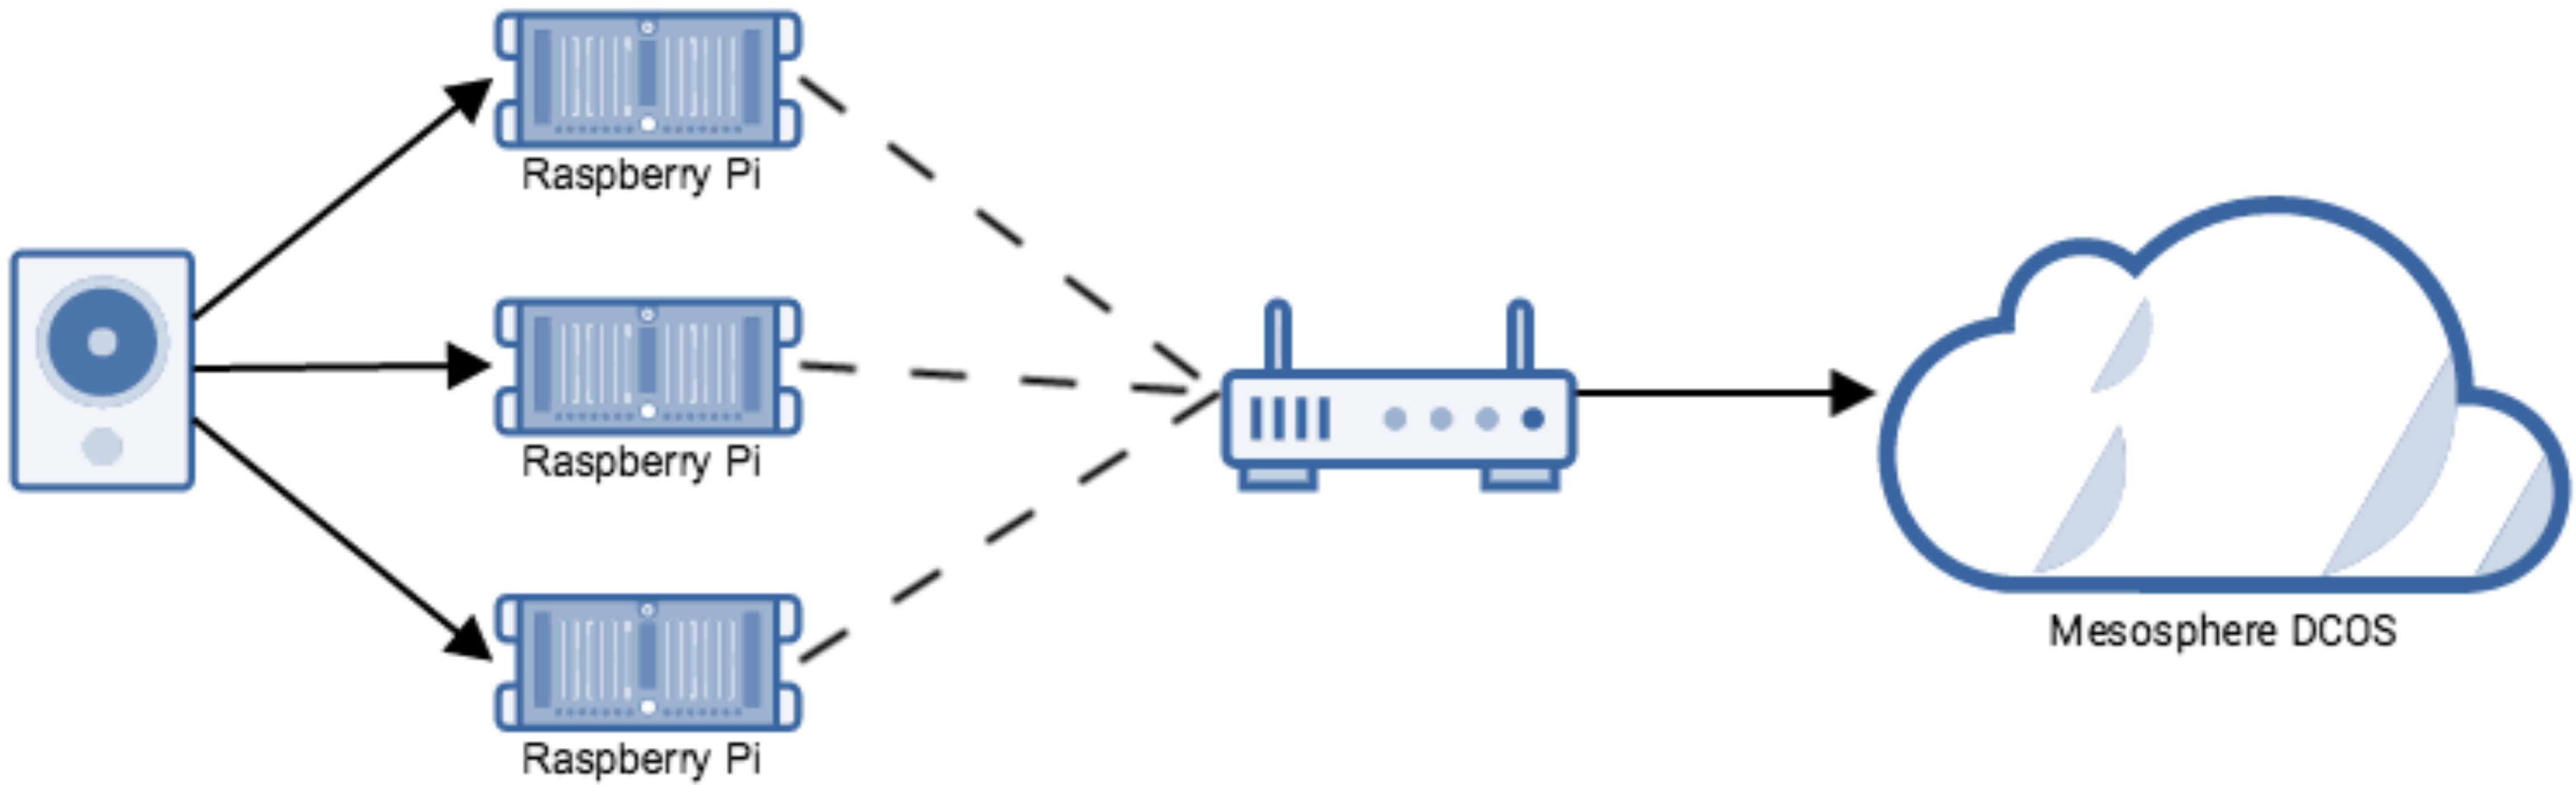
Using Distributed Systems to Ingest Data from the  
**Internet of Pis**



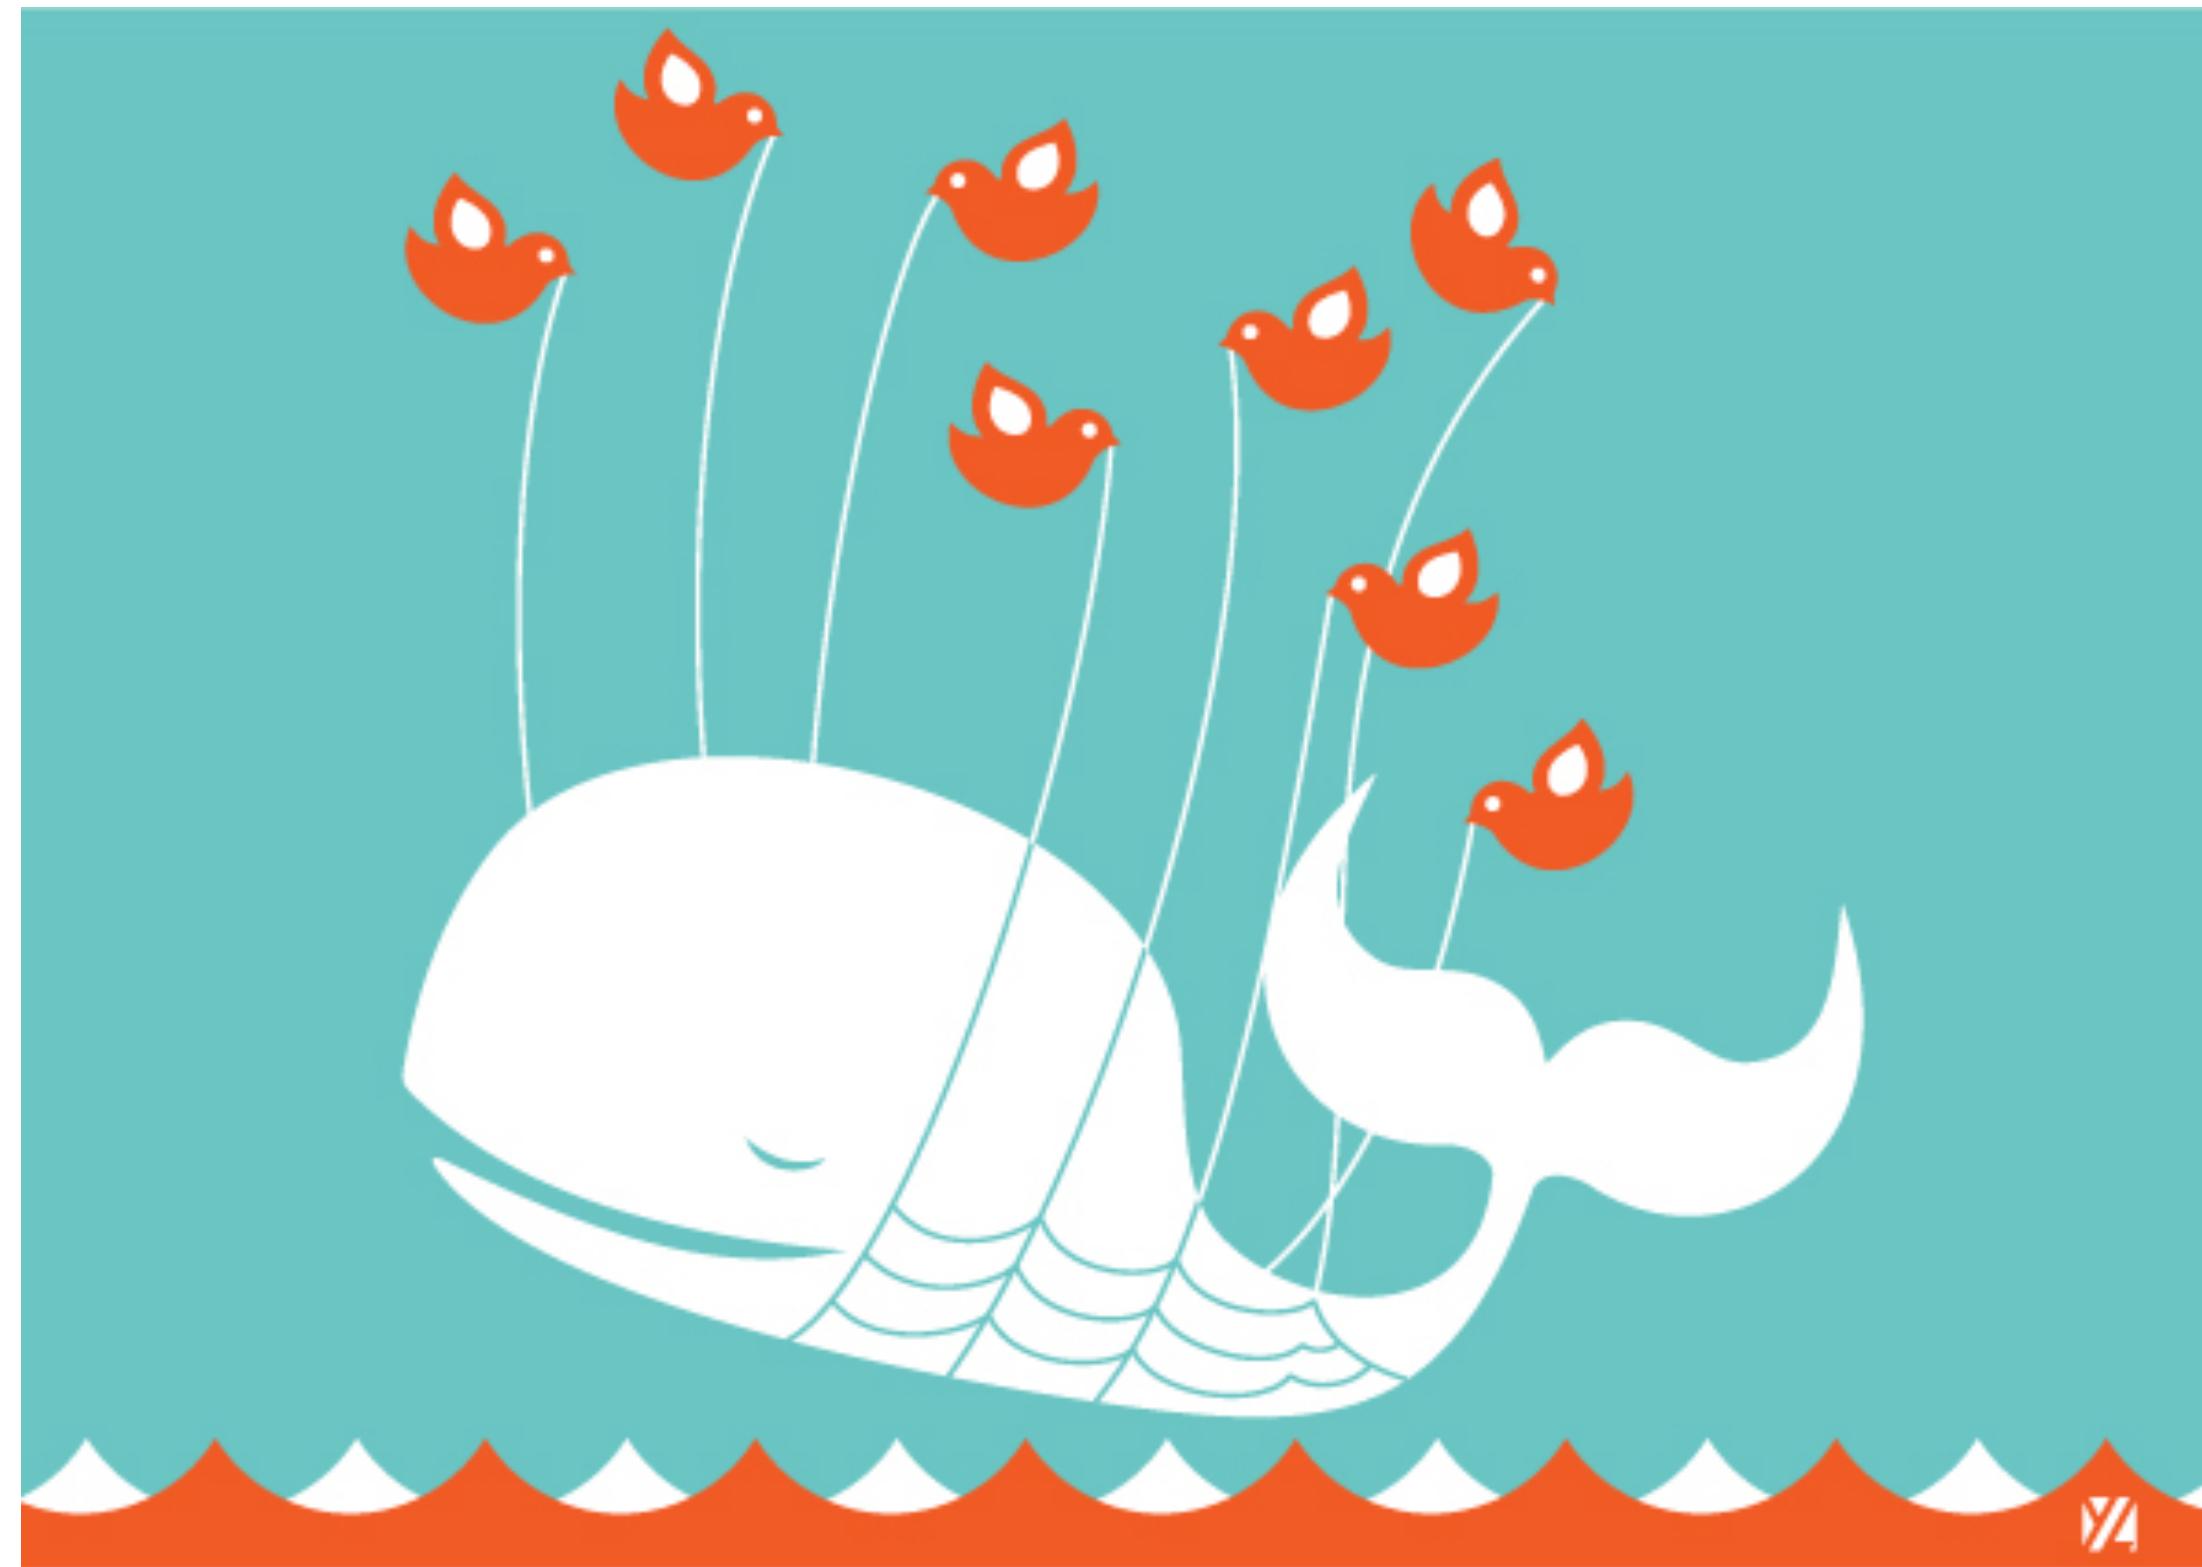
MESOSPHERE

# What's the point?

# Internet of Pis



# Scaling manually is hard!

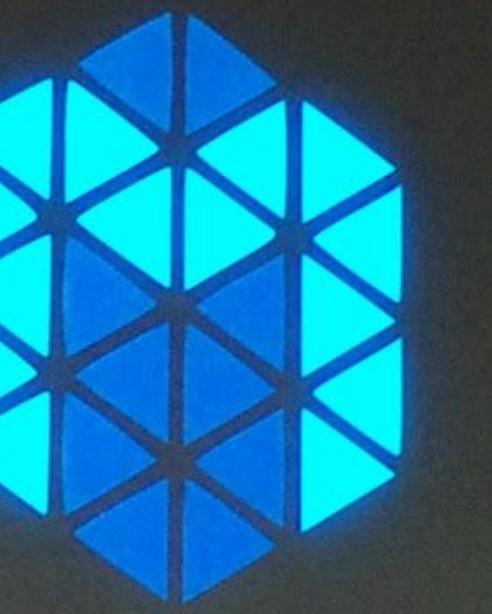


- Mesos is credited with “slaying the fail whale”.
- Twitter’s user growth outstripped their infrastructure growth.
- Mesos, a cluster resource manager, let them scale.

Siri: Today  
Third Generation



powered by





Learn    Product    Documentation    Blog

Get Started

# PUT YOUR DATACENTER ON AUTOPilot.

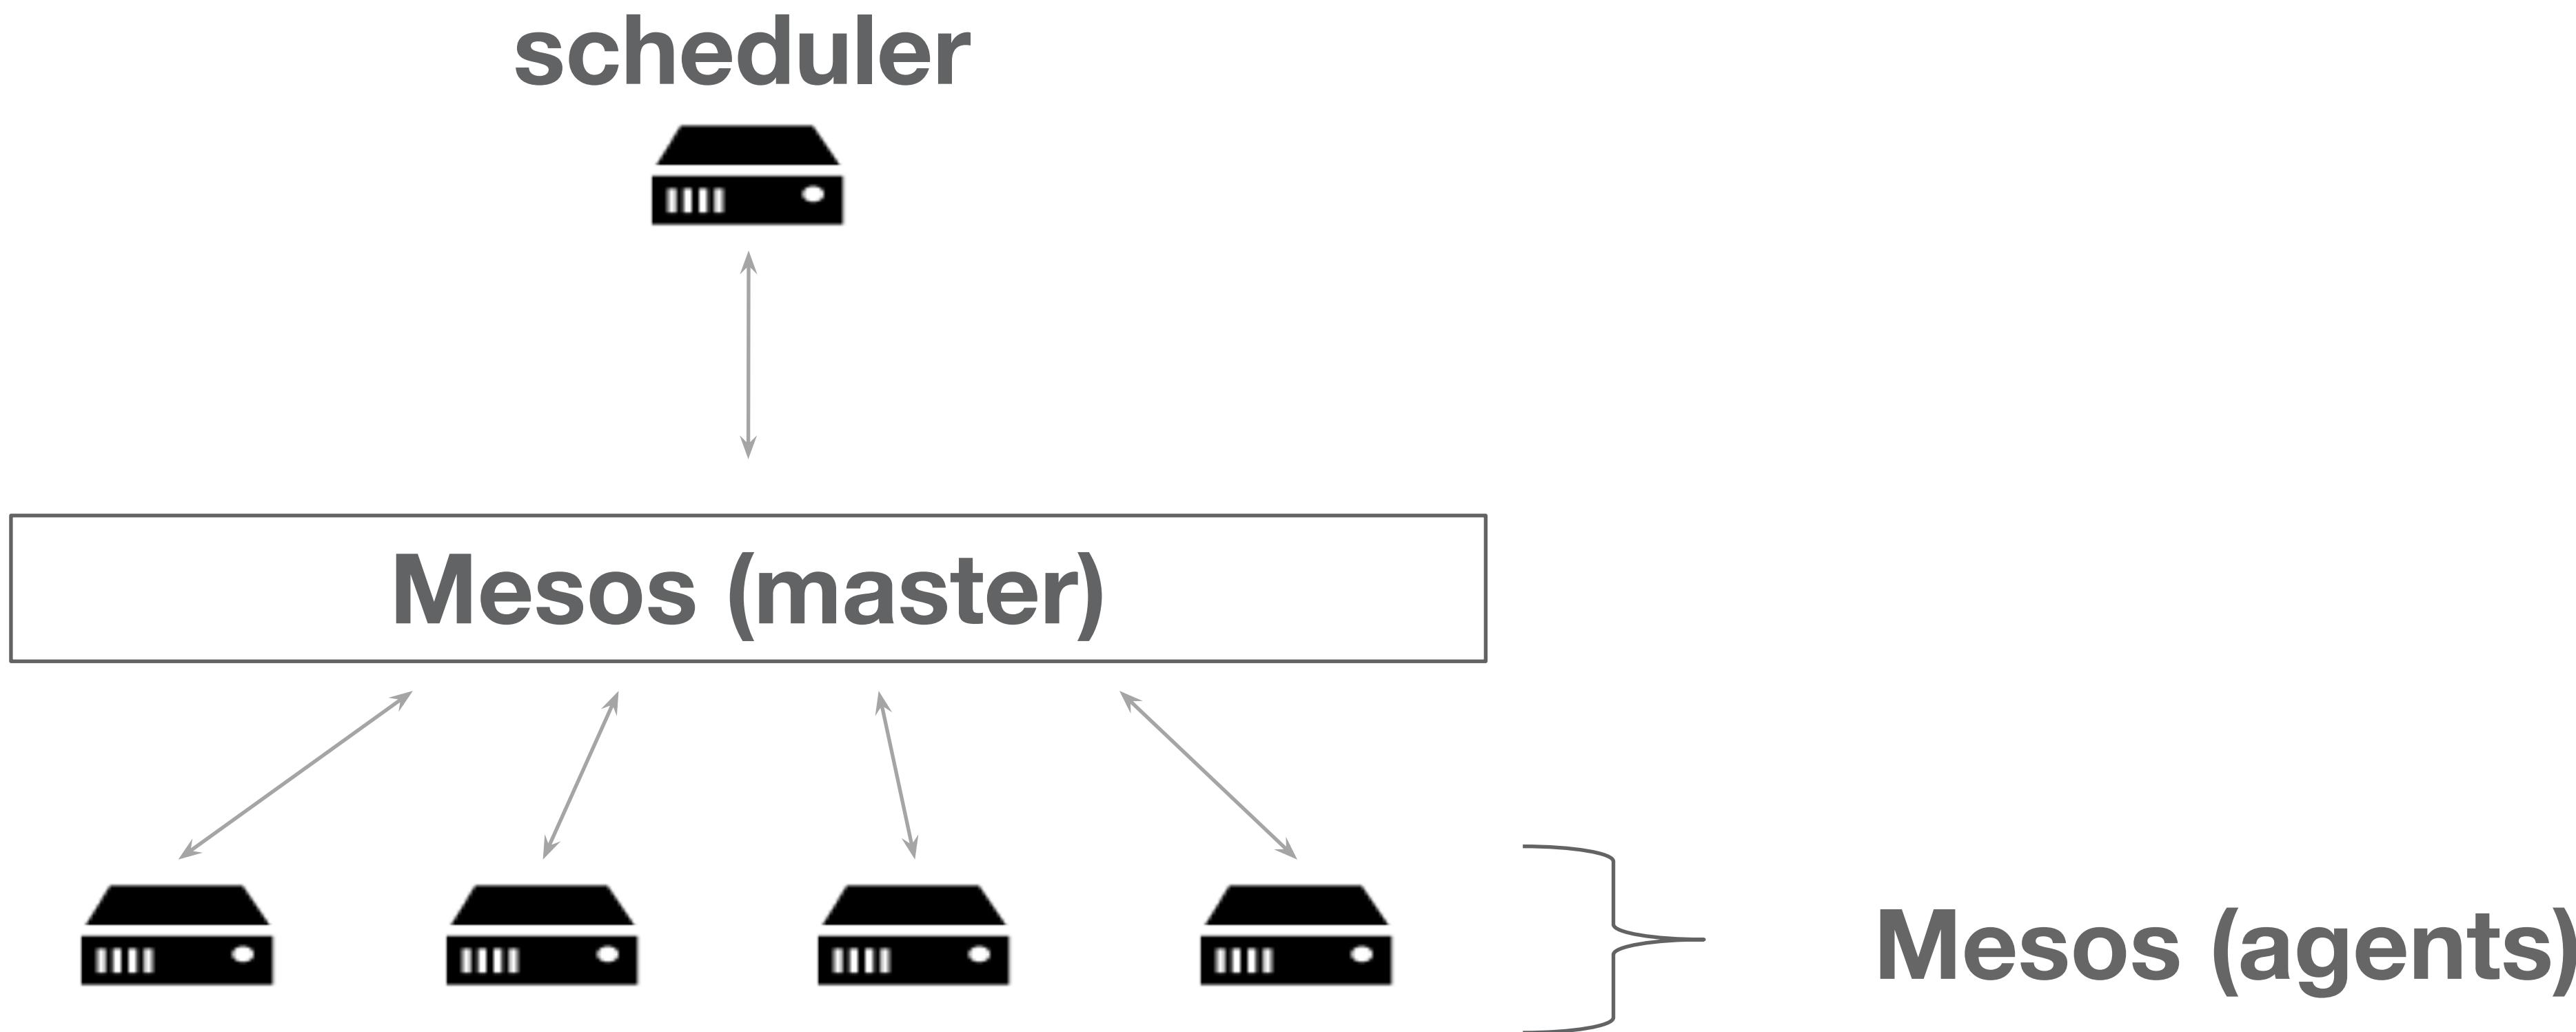
Deploy apps faster. Scale massively. Automate effortlessly.

Get Started Today

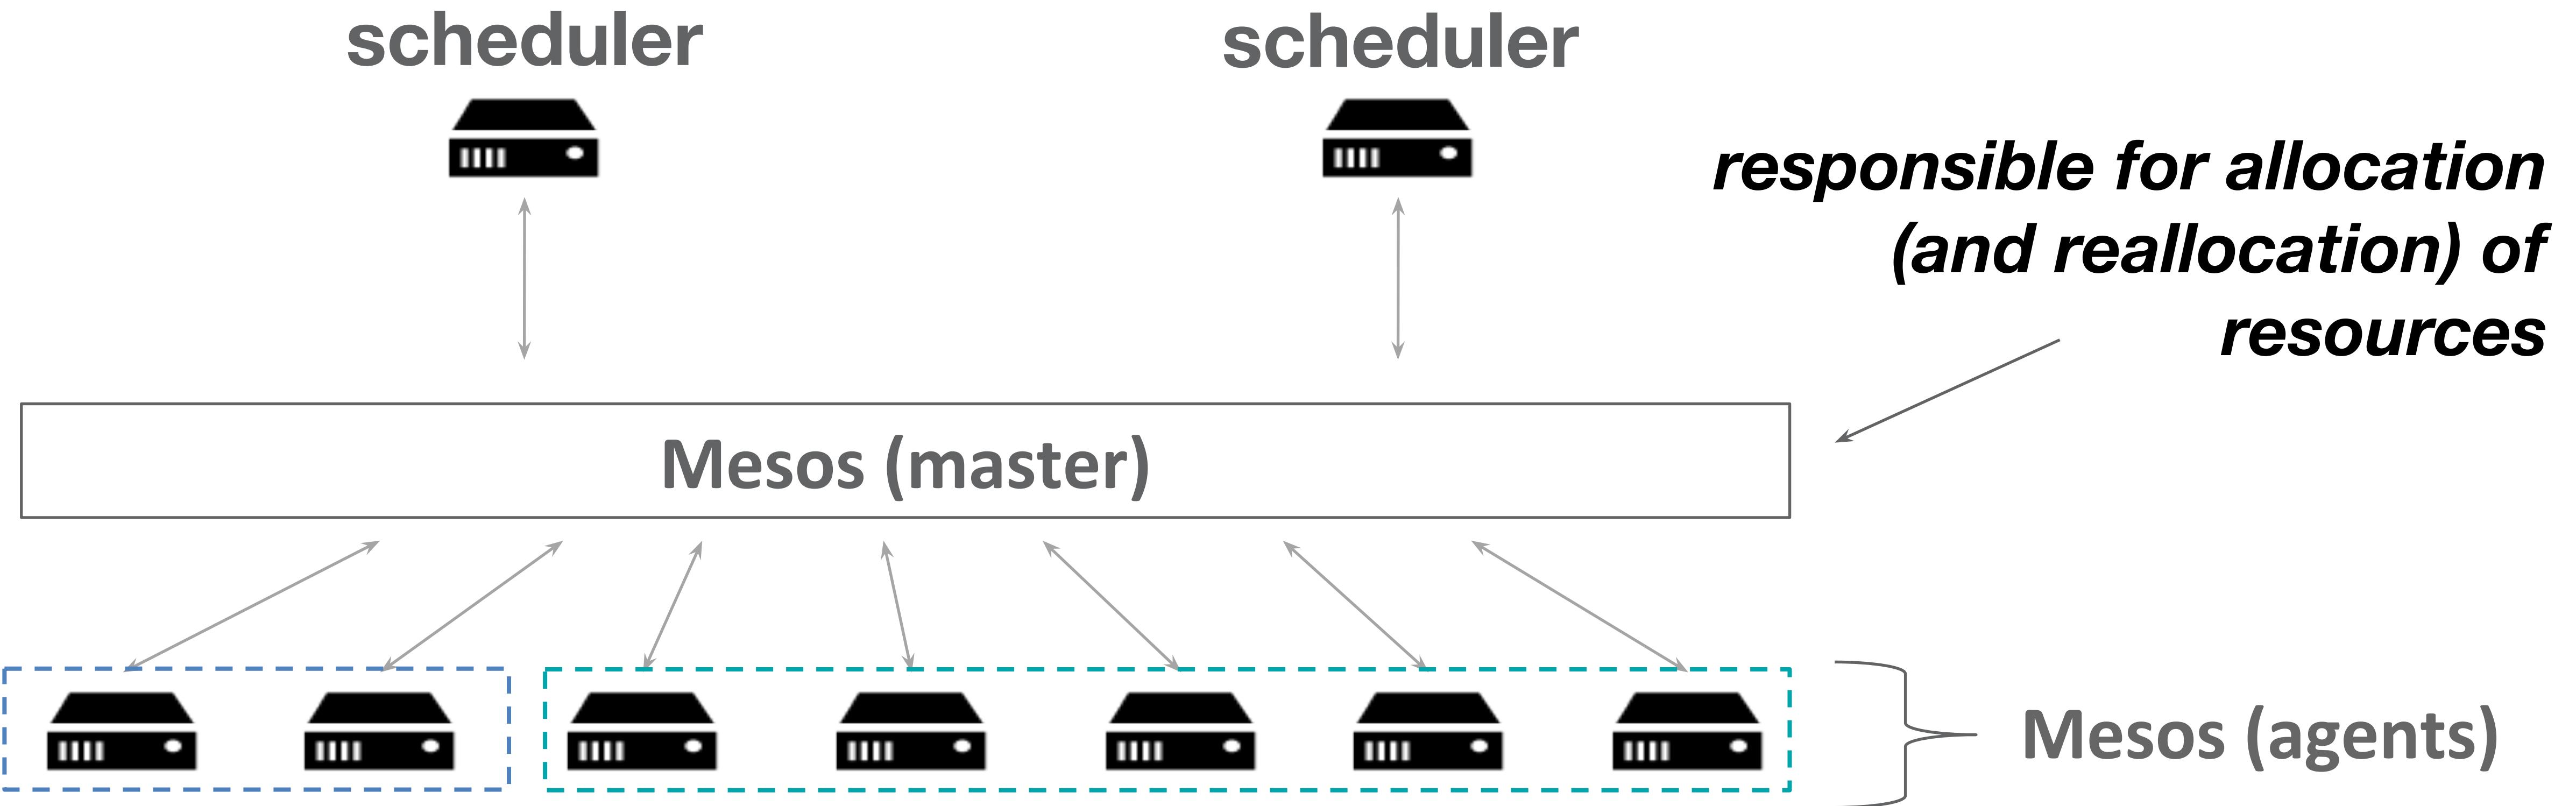
Learn More

# Mesoswho

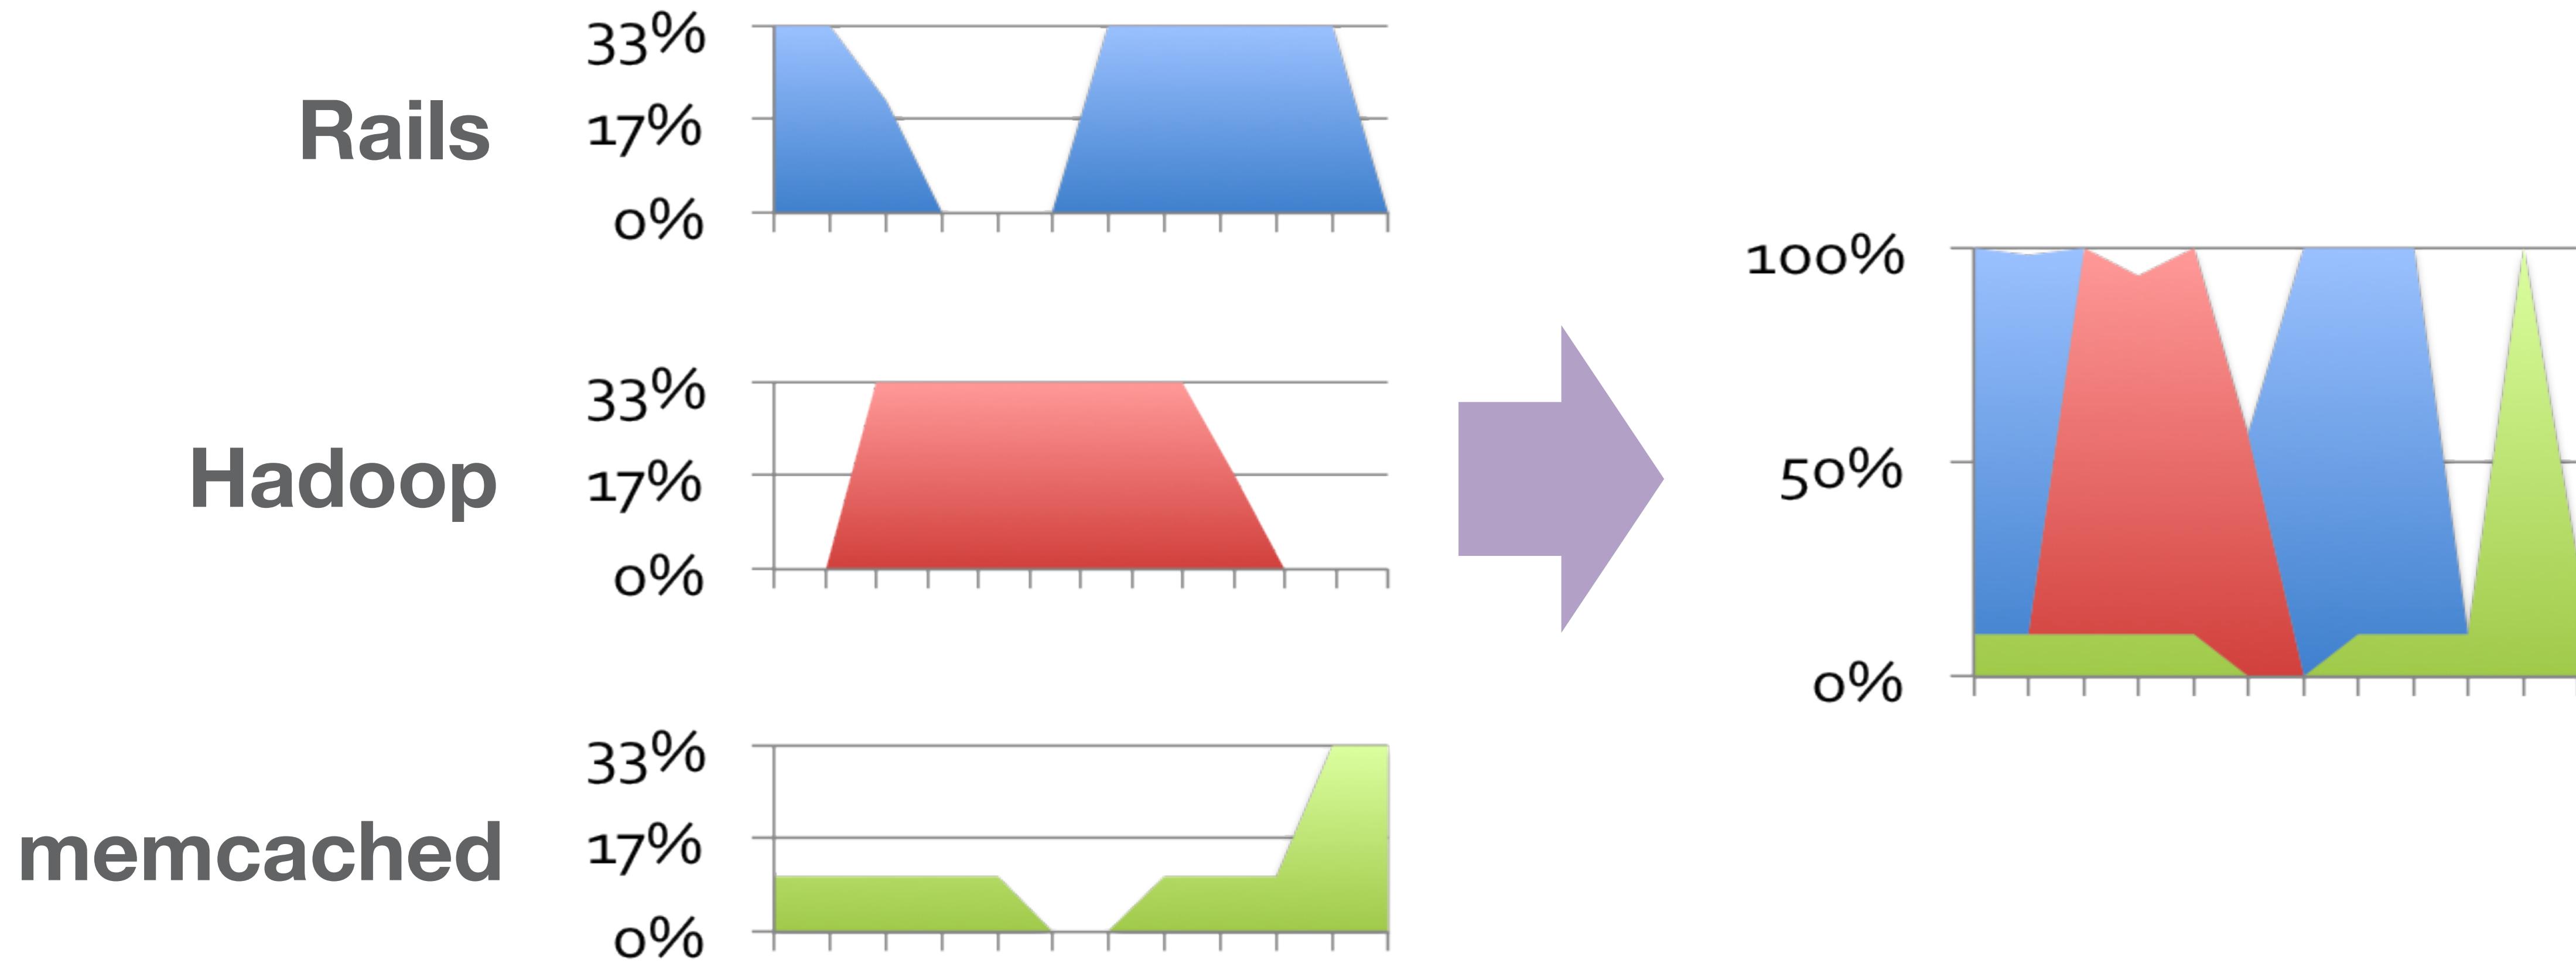
# Mesos: Level of Indirection



# Mesos: Level of Indirection



# Mesos: Improves Utilization



# Mesos: Battle Tested

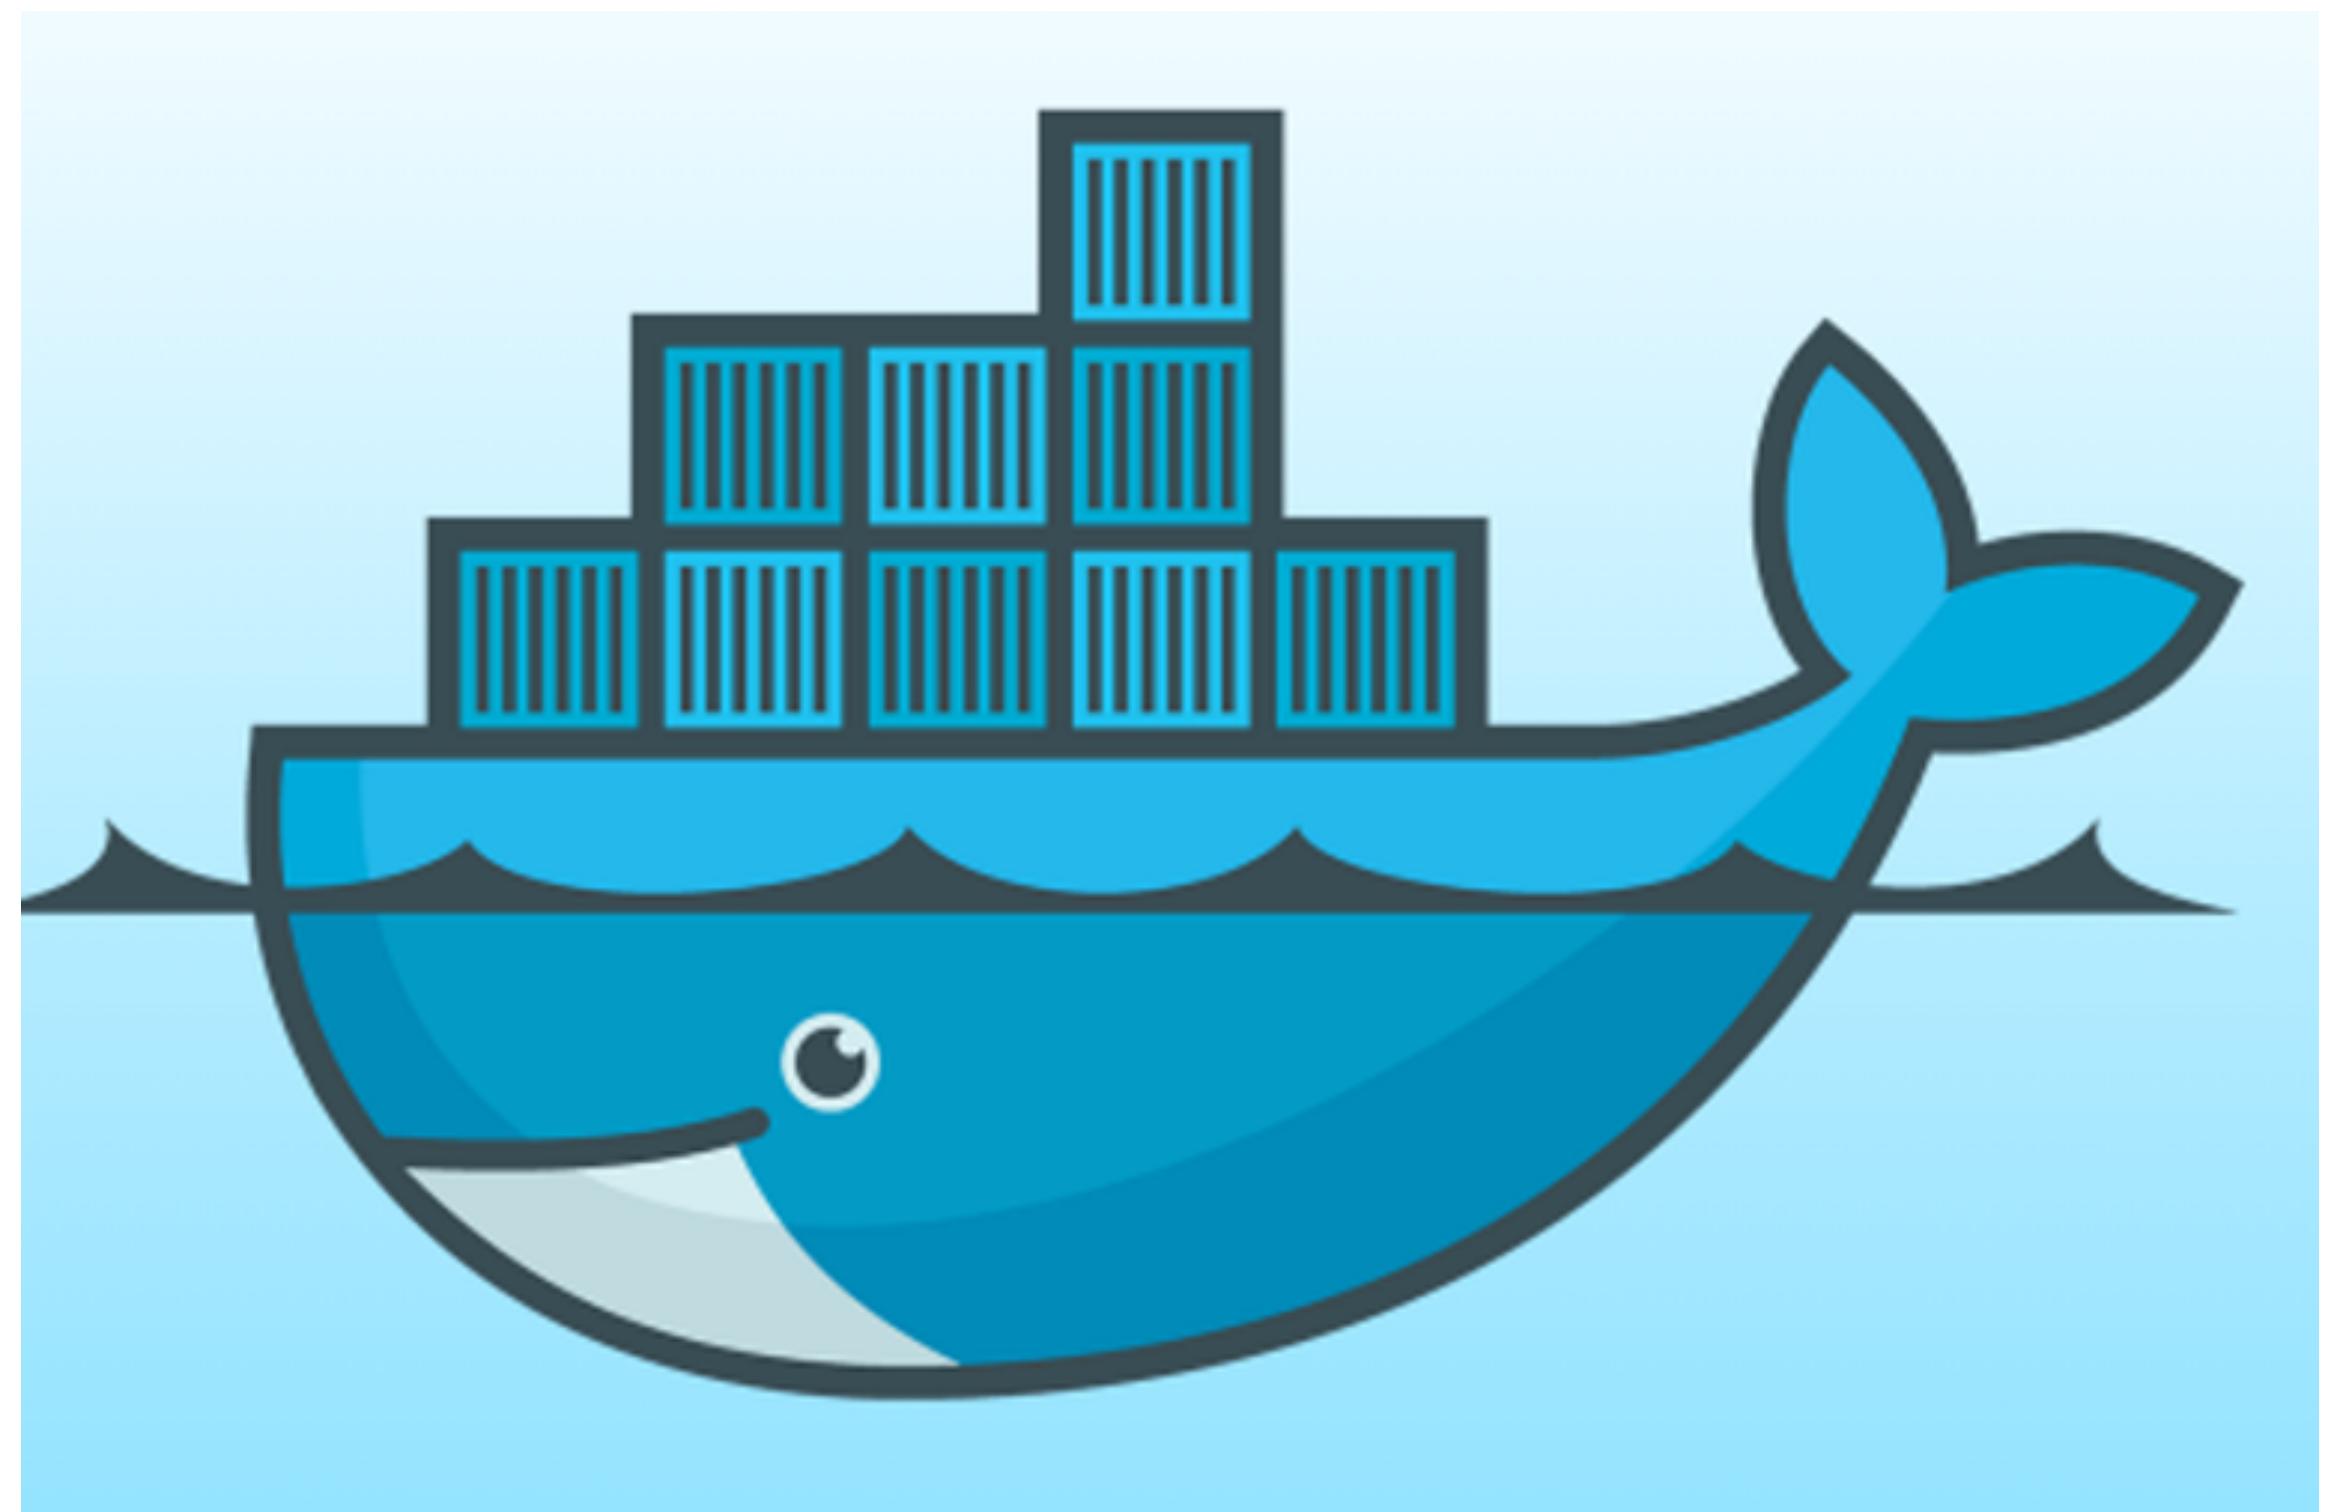
- A top-level Apache project
- A cluster resource negotiator
- Scalable to 10,000s of nodes
- Fault-tolerant, battle-tested
- An SDK for distributed apps



# Containerwhat

# What are containers?

- Shippable unit of execution.
- Take an operating system, bundle together libraries and application files and ship it around.
- Allows deployment environments (i.e. servers) to be homogenous, “Cattle, not pets”



# Workshop!

<http://hackers-at-berkeley.mesosphere.io>

Tyler's

# Philosophy of the Distributed Champion

1. Contain



2. Deploy

3. Discover



# First Steps

# Set up Docker

Docker is a format and tool that allows you to build and run containers.

We'll use it to package our applications for deployment on the cluster.

1. Install Docker on your local machine: <https://docs.docker.com/installation/>
2. Sign up for a Docker Hub account: <https://hub.docker.com>

# Set up the DCOS CLI

DCOS is Mesosphere's Mesos distribution.

We will use the DCOS CLI, a command line client, to easily manage and deploy applications on a DCOS cluster.

1. Install the DCOS CLI: <https://docs.mesosphere.com/install/cli>,  
using the hostname `http://dcos.hackers-at-berkeley.mesosphere.io`

# Get the example files

1. Clone the workshop Github repository to get our example files: <https://github.com/mesosphere/hackers-at-berkeley>

(for git help, try <https://help.github.com/articles/cloning-a-repository/>)

# Exercise 1

Build and run a Python webservice

# Exercise 2

Deploy a Python webservice

# Exercise 3

Create and deploy a Python webservice that talks to Cassandra

# Build & Deploy

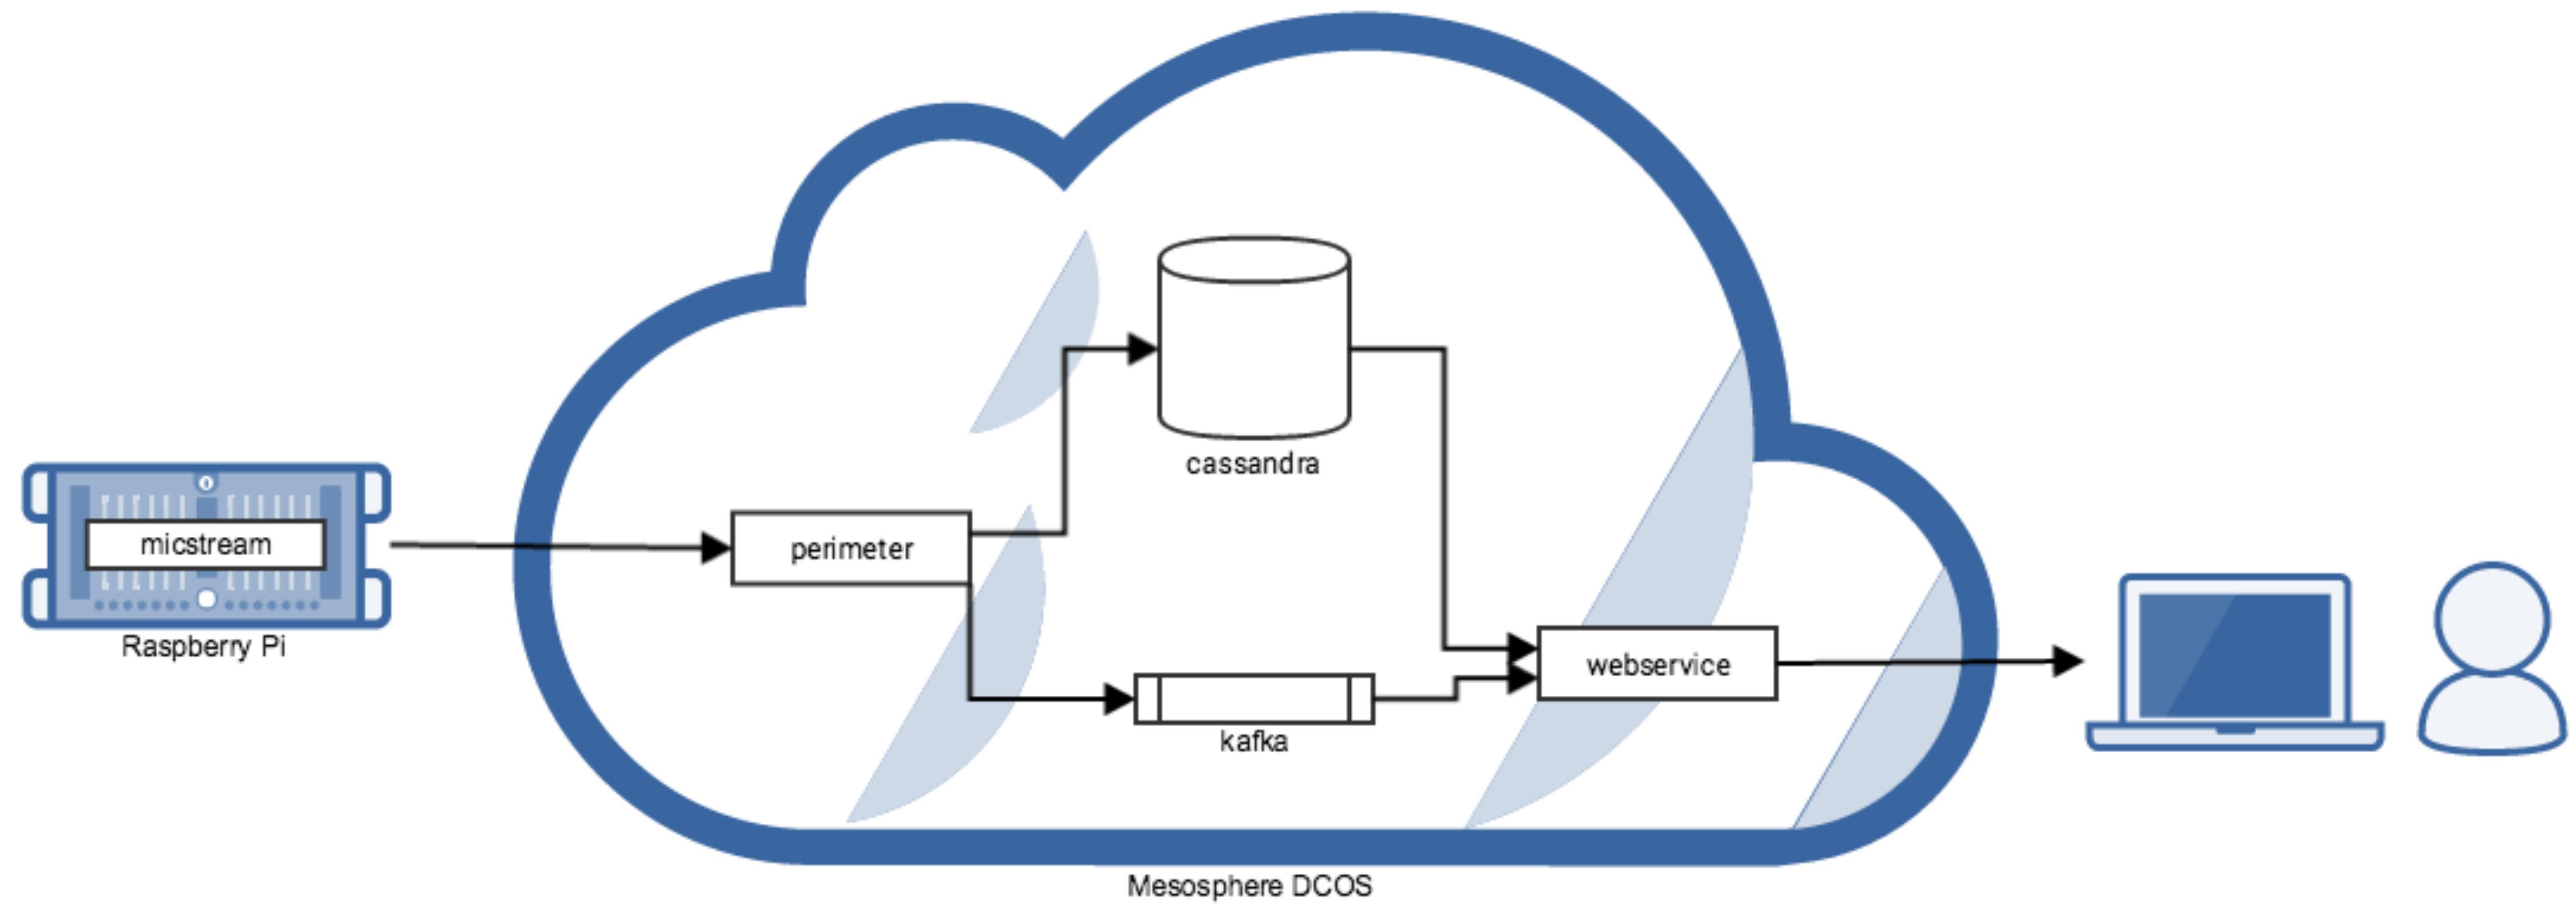
## 1. Build your Docker web service:

```
# from the Github repository root directory:  
cd webservice  
  
# now edit marathon.json, placing your  
# username in the "id" field. then, do:  
docker build -t awesomeuser/my-webapp .
```

## 2. Push the image to the Docker Hub:

```
docker build -t awesomeuser/my-webapp-name
```





# Thanks!

(come and work for us)

# ACT 1 - make something

1. make a WSGI app! (Flask, Django, etc...)
  - examples/snippets/http\_app.py
  - examples/frontend\_service/front\_end.py
2. containerize it!
  - examples/frontend\_service/Dockerfile (note the port)
  - docker build -t <your\_dockerhub\_name>/<app\_name>
  - docker push <your\_dockerhub\_name>/<app\_name>

# ACT 2 - Deploy it

- Marathon is a scheduler that runs processes on Mesos based on your specifications (host types, ports, memory/disk usage, etc...)
- We can tell it to run cli commands, arbitrary executables, Docker containers, etc... using an http POST.
  - examples/frontend\_service/marathon.json
  - `curl -X POST <host>:8080/v2/apps \  
-d @<your_file>.json \  
-H "Content-type: application/json"`

# How to see it:

1. In the marathon UI, check out your app (find it by the name you set in the JSON you POST'd to marathon)
2. You'll see something like “ip-10-0-4-67.us-west-2.compute.internal”
3. find its public IP in our “magic table” :)
  - or [find-hack.mesosphere.io/<app\\_name>](http://find-hack.mesosphere.io/<app_name>)

# ACT 3 - Service Discovery

- Our apps will be scheduled to ports that are available on the machine, which means they cannot be predicted ahead of time.
- We can use SRV DNS records to find both an (internal) IP and a port for a name
  - examples/snippets/marathon\_SRV\_discovery.py
  - `srvlookup.lookup('myapp', 'tcp', 'marathon.mesos')`